

Contextual bandits

David S. Rosenberg

NYU: CDS

October 13, 2021

Contents

1 The contextual bandit problem

2 Policies

The contextual bandit problem

The contextual bandit

A **contextual bandit problem** proceeds in rounds of the following steps:

- ➊ Observe input/context X .
 - ➋ Take action A .
 - ➌ Receive reward $R \in \mathbb{R}$.
- Action A may depends on X and previously observed (X, A, R) triples.

Example: News article recommendation

- Consider a news website.
- Every day there are 10 top stories.
- We want to highlight one for each user.
 - Choice should be **personalized**.
- What is the action?
 - Selecting one of the top 10 stories
- What is the context?
- What is the reward?

The context

- What is the context?
 - Information about the user, if any (e.g. demographics)
 - Geographic location
 - User identifier (we can learn latent features, collaborative filtering style)
- Summary of recent actions
 - Recent reading history
 - Lifetime reading history
 - Shared articles
 - “Friends” (individuals they’ve shared articles with in the past)

- When we give a formal definition of contextual bandits in the next section, we'll see that including a summary of recent actions seems to violate the i.i.d. assumption about the contexts.
- Strictly speaking, including recent actions as part of the context could introduce dependencies between consecutive contexts for the same individual, since an individual is much more likely to have read story 1 by the second time we see the individual if they were recommended story 1 the first time they were seen.
- If the bandit were to be re-started with enough frequency that each individual is usually seen at most once in any run, then these dependencies are unlikely to show up often, and so it seems reasonable to include recent action as part of the context. For example, we might restart the bandit whenever the top 10 news stories changes...

The reward

- What can we use as a reward signal?
- Click (Y/N)
- Spent more than 30 seconds on article page (Y/N)
- More complicated function of time spent reading article
- Was article shared or favorited? (Y/N)
- Figuring out the right reward signal is nontrivial.
 - Requires domain understanding.
 - May need tweaking over time.

Context / reward examples

- User 325. $X = \{\text{Likes sports articles}\}$.
- Actions / rewards
 - Action 1: “Tom Brady retirement” **Reward: 10**
 - Action 2: “Player has meltdown after argument” **Reward: 2**
 - Action 3: “Government considers ban for actor using drugs” **Reward: 3**
- User 823. $x = \{\text{Likes human-interest stories}\}$
- Actions / rewards
 - Action 1: “Tom Brady retirement” **Reward: 1**
 - Action 2: “Player has meltdown after argument” **Reward: 5**
 - Action 3: “Government considers ban for actor using drugs” **Reward: 0**

- Each action is a possible news article (represented here by the title).
- In the terminology of our discussion of causal inference, the reward for each action is a potential outcome.
- We only get to observe the reward corresponding to the action that was taken (or “treatment” given, in the causal inference terminology).
- Terminology note: Some authors refer to the outcomes we don’t observe as “counterfactual” (e.g. [MW15, Ch. 2]).
- Other authors use “counterfactual” to refer to all the potential outcomes that can happen (e.g. [HR20, p. 4]. And one of these counterfactuals, the observed outcome, is also “factual”.
- Some authors are careful to avoid the word “counterfactual” because of this ambiguity.
- Just be aware of the different usages.

The rewards

In each round,

- a reward is generated for each possible action $a \in \mathcal{A} = \{1, \dots, k\}$.
- These rewards are conditioned on the context $X \in \mathcal{X}$.
- We'll represent the k rewards by a **reward vector**

$$R = (R(1), \dots, R(k)) \in \mathbb{R}^k.$$

- Only a single entry of R is revealed, namely $R(A)$, where A is the action played for that round.

Probabilistic model for contextual bandit

- Context and reward vector are related:
 - The same action will get different rewards in different contexts.

Stochastic contextual k -armed bandit model

- Context and reward vector $(X, R) \in \mathcal{X} \times \mathbb{R}^k$ **drawn jointly** from P .
- Context and reward pairs are i.i.d. over time:

$$(X, R), (X_1, R_1), \dots, (X_t, R_t) \text{ i.i.d. } \sim P.$$

- For contextual bandits to add something interesting to our original bandit formulation, there needs to be some dependence between the rewards received for each action and the context.

Action selection

- Action at round t is A_t .
- At beginning of round t , the **history**, or previous **observation sequence** is

$$\mathcal{D}_t = \left((X_1, A_1, R_1(A_1)), \dots, (X_{t-1}, A_{t-1}, R_{t-1}(A_{t-1})) \right).$$

- At round t , action A_t may depend on context X_t and history \mathcal{D}_t .
- Note that we **cannot** say $A_t \perp\!\!\!\perp R_t$ – why? (note R_t not $R_t(A_t)$)
- Because A_t depends on X_t , and R_t depends on X_t .
 - Information about R_t can propagate to A_t through X_t .

Action and reward are **conditionally** independent given context

We can say that $A_t \perp\!\!\!\perp R_t \mid X_t$ for each t .

- Note that $A \perp\!\!\!\perp R \mid X$ is the exact counterpart to the “ignorability” assumption in causal inference: $(Y(0), Y(1)) \perp\!\!\!\perp W \mid X$. The reward vector $R = (R(1), \dots, R(k)) \in \mathbb{R}^k$ corresponds to the potential outcome vector $(Y(0), Y(1)) \in \mathbb{R}^2$. The action $A \in \mathcal{A}$ corresponds to the treatment indicator $W \in \{0, 1\}$, and the covariate $X \in \mathcal{X}$ has the same interpretation in each setting.

Stochastic k -armed contextual bandit

Stochastic k -armed contextual bandit

- 1 Environment samples **context** and **reward vector** jointly, iid, for each round:

$$(X, R), (X_1, R_1), \dots, (X_T, R_T) \in \mathcal{X} \times \mathbb{R}^k \text{ i.i.d. from } P,$$

where $R_t = (R_t(1), \dots, R_t(k)) \in \mathbb{R}^k$.

- 2 For $t = 1, \dots, T$,

- 1 Our algorithm **selects action**/arm $A_t \in \{1, \dots, k\}$ based on X_t and history

$$\mathcal{D}_t = \left((X_1, A_1, R_1(A_1)), \dots, (X_{t-1}, A_{t-1}, R_{t-1}(A_{t-1})) \right).$$

- 2 Our algorithm **receives reward** $R_t(A_t)$.

- We **never observe** $R_t(a)$ for $a \neq A_t$.

- It might look cleaner to say that at the beginning of every round, the environment generates $(X_t, R_t) \in \mathcal{X} \times \mathbb{R}^k$ from P . But we want to be very clear that $(X_1, R_1), \dots, (X_T, R_T)$ are
 1. generated i.i.d. and are
 2. generated before any of the actions A_1, \dots, A_T are generated.

Policies

- Policies give some structure to action selection.
- A policy at round t
 - gives a conditional distribution over the action A_t to be taken
 - conditioned on the current context X_t and the history \mathcal{D}_t .
- We'll denote the policy at round t as $\pi_t(\cdot | X_t, \mathcal{D}_t)$.
- Choosing an action according to policy π_t means we choose A_t **randomly** s.t.

$$\mathbb{P}(A_t = a) = \pi_t(a | X_t, \mathcal{D}_t).$$

Optimal policy

- Suppose we knew the function

$$r(x, a) = \mathbb{E}[R \mid A = a, X = x],$$

which gives the expected reward for taking action a in context x .

- Then the optimal policy would be

$$\pi_t^*(a \mid X_t, \mathcal{D}_t) = \mathbb{1} \left[a = \arg \max_a r(X_t, a) \right].$$

- Of course, we don't know $r(x, a)$, but can we estimate it?

Example: “direct method”

- We don't know $r(x, a)$, but we can use \mathcal{D}_t as training data:

$$\left(\underbrace{(X_1, A_1)}_{\text{input}}, \underbrace{R_1(A_1)}_{\text{label / response}} \right), \dots, \left(\underbrace{(X_{t-1}, A_{t-1})}_{\text{input}}, \underbrace{R_{t-1}(A_{t-1})}_{\text{label / response}} \right).$$

- Estimating $r(x, a)$ is a regression problem!
- Let $\hat{r}_t(x, a) = \text{TrainingAlgorithm}(\mathcal{D}_t)$.
- The policy for the **direct method** is defined as

$$\pi_t(a \mid X_t, \mathcal{D}_t) := \mathbb{1} \left[a = \arg \max_a \hat{r}_t(x, a) \right].$$

- This is a **pure exploitation** method.

Some other approaches

- ϵ -greedy is an obvious extension of the direct method.
- Thompson sampling: prior is over models $\hat{r}_t(x, a)$
 - equivalently, prior is over model parameters
- Policy gradient: directly optimizing over the policy to improve expected reward
 - we'll return to this in a few weeks as a warm-up for REINFORCE.

References

- The term contextual bandit was introduced in [LZ07], but the idea has been around much longer.
- A nice history of contextual bandits is given in [TM17], which cites a 1979 paper as the first appearance of contextual bandits.

References I

- [HR20] Miguel A. Hernán and James M. Robins, *Causal inference: What if*, Boca Raton: Chapman & Hall/CRC, 2020,
<https://www.hsph.harvard.edu/miguel-hernan/causal-inference-book/>.
- [LZ07] John Langford and Tong Zhang, *The epoch-greedy algorithm for contextual multi-armed bandits*, Proceedings of the 20th International Conference on Neural Information Processing Systems (Red Hook, NY, USA), NIPS'07, Curran Associates Inc., 2007, pp. 817–824.
- [MW15] Stephen L. Morgan and Christopher Winship, *Counterfactuals and causal inference*, 2 ed., Cambridge University Press, 2015.
- [TM17] Ambuj Tewari and Susan A. Murphy, *From ads to interventions: Contextual bandits in mobile health*, Mobile Health, pp. 495–517, Springer International Publishing, 2017.